

Chapter 1

Implementation of ALGT

In this chapter, the codebase of ALGT is explored. All important technical choices are given, making a reimplementaion of core-ALGT possible: the representation of the syntax, functions and natural deduction rules; together with their usage: the parser and interpreters for functions and rules. The algorithms are accompanied with Haskell-snippets or pseudocode, so some familiarity with basic haskell is required for this chapter. These illustrate the algorithms, but are often simplified. Complications for additional features, often conceptually simple yet tremendously practical, are omitted.

This chapter does not cover the abstract interpretation; these algorithms are already explained in detail in chapter ??.

1.1 Representation and parsing of arbitrary syntax

The first aspect of any programming language is its syntax. ALGT allows to denote these explicitly, on which a parser can be based as can be seen in ??. Of course, the BNF has a representation within ALGT, which is given by the following construct:

```
1 -- Representation of a single BNF-expression
2 data BNF      = Literal String      -- Literally parse 'String'
3               | BNFRuleCall Name    -- Parse the rule with the given name.
4               | BNFSeq [BNF]        -- Sequence of parts
```

The most fundamental element here is the `Literal`, which is the terminal given in the string. Calling a non-terminal or builtin value is represented with `BNFRuleCall`. By using external definitions for the builtins, the main representation can be kept small and clean. At last, single elements can be glued together using `BNFSeq`.

All these expressions are bundled into a syntax:

```
1 {-Represents a syntax: the name of the rule + possible BNFS -}
2 data Syntax    = BNFRules
3               { bnf          :: Map TypeName [BNF]
4               , lattice      :: Lattice TypeName
5               }
```

All the syntactic forms are saved into the dictionary `bnf`, which contains a mapping from the name of the non-terminal onto all possible choices for that syntactic form. As a syntactic form also is a type for the metafunctions, this dictionary also doubles as store for known types.

The second responsibility of the syntax is keeping track of the supertyping-relationship. The lattice-data structure keeps track of what type is a subtype of what other types and will play a major role in the typechecker.

1.1.1 Target language parsing

A syntax as above can be interpreted as a program, acting on an input string generating a parse-tree. The parser, which uses the *Parsec*-package in the background, is constructed recursively as can be seen in figure 1.1.

The entry point for the parser is `parseNonTerminal`, which takes the name of the rule that should be parsed - together with the syntax itself. In this syntax, the relevant choices are searched. These choices are tried one by one in `parseChoices`. The actual parsing in `parseChoices` is delegated to `parsePart`, which does the actual interpretation: tokens (`Literal`) are parsed literally, sequences (`BNFSeq`) are parsed using `parsePart` recursively. If a non-terminal (`RuleCall`) is encountered, `parsePart` calls `parseNonTerminal`, closing the loop.

1.2 Target program representation

Target programs are represented as *parsetrees*. The data structure responsible is structured as following:

```

1 data ParseTreeA
2   = MLiteral      {ptaContents :: String}
3   | MInt          {ptaInt :: Int}
4   | PtSeq         {ptaPts :: [ParseTreeA]}
```

This implementation falls apart in the concrete values (`MLiteral` for strings and `MInt` for numbers) and a node to combine parts into longer sequences: `PtSeq`. `PtSeq` acts as node element in the parsetree, whereas the concrete values are the branches.

In the actual implementation more information is tracked in the parsetree, such as what syntactic form constructed the parsetree, the starting position of each token and its length. For clarity, this extra information is omitted in this text.

The string `(1 + 2)` -when parsed with `e` from STFL- is represented by the parsetree `PtSeq [MLiteral "(", PtSeq [MLiteral "1", MLiteral "+", MLiteral "2"], MLiteral ")"]`. Its graphical representation can be seen in figure 1.2

1.3 Metafunctions

Just as the syntax, the metafunctions are explicitly represented within ALGT. Remember that all metafunctions operate on a parsetree -a part of the target program. The representation of functions is thus closely linked with these parsetrees, but yet straightforward:


- The `Function` data structure contains the body of the function -multiple `clauses`- and the type of the function (the types of the arguments and the type it'll return). These types are given explicitly in the source code, thus no inference is needed.
- A `Clause` contains the `Expression` returned and zero or more expressions which will performing the pattern matches.
- An `Expression` acts both as the pattern match representation and as the expression constructing a new parsetree as result of the function.

```

1
2 -- Seaches the given rule in the syntax, tries to parse it
3 parseNonTerminal      :: Syntax -> Name -> Parser ParseTree
4 parseNonTerminal syntax@(BNFRules syntForms _) ruleToParse wsModeParent
5   | ruleToParse `M.notMember` syntForms
6     = fail $ ruleToParse++" is not defined in the given syntax"
7   | otherwise         = do   let choices      = syntForms ! ruleToParse
8                             parseChoice syntax ruleToParse choices
9
10
11
12
13 -- Search one of the choices that can be parsed
14 parseChoices          :: Syntax -> Name -> [BNF] -> Parser ParseTreeLi
15 parseChoices _ ruleToParse _ []
16   = fail $ "All choices depleted, can not parse anything"
17   = fail $ "Could not parse expression of the form "++ruleToParse
18 parseChoices syntax ruleToParse (choice:rest)
19   = try (parsePart syntax choice)
20     <|> parseChoices syntax name rest
21
22
23
24
25 -- Parses a single BNF-part
26 parsePart             :: Syntax -> BNF -> Parser ParseTree
27
28 -- Parse exactly the literal str
29 parsePart _ tp _ (Literal str)
30   = do   string str
31         return $ MLiteral str
32
33 -- Parse a rulecall...
34 parsePart syntax _ bnf@(BNFRuleCall ruleToParse)
35   | isBuiltin bnf
36     = do   let parser          = getParserForBuiltin bnf
37           parsedStr          <- parser
38           return $ MLiteral parsedStr
39   | otherwise
40     = parseNonTerminal rules ruleToParse wsMode
41
42 -- degenerate case for sequences: a sequence with a single element
43 parsePart syntax (BNFSeq [bnf])
44   = parsePart syntax tp bnf
45
46 -- parse all elements of the sequence in order
47 parsePart syntax (BNFSeq (bnf:bnfs))
48   = do   head    <- parsePart syntax tp bnf
49         tail    <- bnfs |> parsePart' syntax tp
50         return $ PtSeq (head:tail)

```

Figure 1.1: The recursively constructed parser. `MLiteral` and `PtSeq` are part of the target program representation and explained in 1.2

Figure 1.2: Parsetree of $(1 + 2)$

For each usage of an `Expression` (as detailed in section 1.3), an element is declared in the algebraic data type `Expression`:

- `MVar` is used to represent metavariables.
- `MParseTree` is used when a literal piece of target program is needed.
- `MCall` is a call to another function in scope.
- `MSeq` is a node in the meta-expression, grouping multiple parts to a bigger expression.
- `MAscription` will coerce the embedded expression to be of the given type.
- `MEvalContext` will, when used as pattern, create a hole within `Name` of the form `Expression`. When `MEvalContext` is used to construct a parsetree, it will plug the hole with the embedded expression.

Each `Expression` carries a `TypeName` as first argument, indicating the type of the expression. These types are inferred once the file is parsed.

```

1  -- A single pattern or expression
2  data Expression
3      = -- a variable
4          MVar TypeName Name
5          -- a 'value'
6          | MParseTree TypeName ParseTree
7          -- function call; not allowed in pattern matching
8          | MCall TypeName Name Builtin [Expression]
9          -- A node, containing multiple expressions
10         | MSeq MInfo [Expression]
11         -- checks wether the expression is built by this smaller type.
12         | MAscription TypeName Expression
13         -- describes a pattern that searches a context
14         | MEvalContext TypeName Name Expression
15
16
17  -- A single clause of the function body
18  data Clause
19      = MClause [Expression] Expression
20
21  -- A full function
22  data Function
23      = Function [TypeName] TypeName [MClause]

```

1.3.1 Typechecker of expressions

All expressions and patterns are typechecked, as type errors are easily made. Forcing parsetrees to be well-formed prevents the creation of strings which are not part of the language, what would result in hard to track bugs later on. Here, an overview of the inner workings of this typechecker (or more precisely `_type` annotator) are given. The pseudocode annotating expressions can be found in 1.5

These internals are simplified, as a type expectation is always available: expressions and patterns are always typed explicitly, as the type signature of the function always gives a hint of what syntactic form^[syntacticType] a pattern or expression is. A type for a pattern indicates what type the pattern should deconstruct, or analogously for expressions, what syntactic form a parsetree would be if the expression was used to construct one. The natural deduction rules, which will be introduced in the following part, have the same typing available and can thus be typechecked with the same algorithm.

As expressions and patterns are **duals** in function of semantics, but the same in syntax and internal representation, the same typechecking algorithm can be used for patterns and expressions. However, some fundamental differences exist between in usage between patterns and expressions. The most striking example are variables: in a pattern context, an unknown variable occurrence is a declaration; in an expression context, an unknown variable is an error. In order to keep the typechecker uniform, the typechecker merely **annotates** types to each part of the expression; checks for unknown variables are done afterwards by walking the expression again.

To type **function calls**, a store γ containing all function signatures is provided. This dictionary γ is built before any typechecking happens by assuming the given function signatures are correct. A store for variables is not necessary, as variable typings are compared after the actual type annotation of expressions: a variable table is constructed, in which conflicts can be easily spotted.

With these preliminaries, we present the actual typechecking algorithm used in ALGT. The algorithm has a number of cases, depending on the kind of expression that should be typed; composite expressions are handled by recursively typing the parts before handling the whole.

Variables

Variables are simply annotated with the expected type. One special case is when two (or more) patterns assign the same variable, such as the clause $f(x, x) = \dots$. This is perfectly valid ALGT, as this clause will match when both arguments are identical. With the type signature $f : a \rightarrow b \rightarrow c$ given, type of x can be deduced even more accurately: the biggest common subtype of a and b , as x should be both an a and a b .

Actual type errors are checked after the initial step of type annotation, when all typing information is already available and inconsistencies can be easily detected.

To catch these inconsistencies between assignment and usage, the following strategy is used:

- First, pattern assignments are calculated; this is done by walking each pattern individually, noting which variables are assigned what types.
- When these individual pattern assignments are known, they are merged. Merging consists of building a bigger dictionary, containing all assignments. If two patterns assign the same variable, compatibility of the types is checked by taking the intersection of both types. If that intersection exists (a single common subtype), then some parsetrees exists which might match the pattern and this common subtype is taken as the type of the variable. If no such subtype exists, a type error is generated.

- With this store of all variable typings at hand, the expression can be checked for *undeclared* variables. This is simply done by getting the assignments of the expression -the same operation as on patterns- and checking that each variable of the expression occurs in the assignment of the patterns. If not, an unknown variable error is issued.
- The last step checks for inconsistencies between declaration and usage, which checks that a variable always fits its use, thus that no variable is used where a smaller type is expected.

The merging algorithm is listed in figure 1.6.

For example, the first clause of `f` (as seen in figure 1.3) has the typing assignments $\{ "x" \rightarrow \{ \text{bool}, \text{int} \} \}$. As no intersection exists between `bool` and `int`, an error message is given.

```

1 not      : bool -> bool
2 not("True")    = "False"
3 not("False")   = "True"
4
5 f           : int -> bool -> ...
6 f(x, x)      = ...

```

Figure 1.3: Example of conflicting variable usage: `x` is used both as `bool` and `int`.

Sequences and string literals

As all expressions have an explicit type expectation, typechecking expressions becomes easier. The possible sequences are gathered from the syntax definition and are aligned against the sequence to be typechecked. Then, each element is compared independently: literals as given in the syntax definition should occur at the same position in the sequence, non-terminals should match the respective subexpressions, as can be seen in figure 1.4.

"1"	"+"	x	Should be typed as	expr
eL	"+"	e		
eL	"::"	type		Literals don't match
eL	e			Not enough elements
eL				Not enough elements

Figure 1.4: An example sequence and possible typings

Functions

Functions are typed using the store containing all the function signatures. As all functions are explicitly typed, it is already available.

First, all the arguments are typed individually; then the return type of the function is compared against the expected type of the pattern/expression. The comparison used is, again, subtyping, as this always gives a sound result:

- When the function is used in an expression, a smaller type will fit nicely in the parsetree.
- When used as an pattern, the function is calculated and compared against the input parsetree. Here, the only requirement is that there exist *some* parsetrees that are common to the argument type and the result type. In this case, the only check should be that some intersection of both types exists. As $expectedtype <: funtiontype$ guarantees this, it is a sufficient condition. While this check is a little *to strict*, it is sufficient for practical use.

Type annotations

Type annotations are used for two means. First, it allows easy capturing of syntactic forms (e.g. `isBool((_:bool)) = "True"`). Secondly, they allow disambiguation of definitions in more complicated grammars.

A type annotations such as `x:T` is typechecked in two steps:

- The first check is that an expression of type `T` can occur at that location. This is easily checked: `T` should be a subtype of the expected type.
- The second check is that `x` can be typed as a `T`. This is done by running the type annotator recursively on `x`, with `T` as expectation.

Evaluation contexts

Evaluation contexts implement searching behaviour: when a parsetree is matched over `e[x]`, a subtree matching `x` is searched within the tree. If no such tree is found; the match fails. When this match is found, both `x` and `e` are available as variables.

The expression in the hole can be some advanced expression that should match the subtree. An other expression can be used in turn to construct a slightly different parsetree;

The explicit typing makes it possible to easily tag `e`, as its type `T` will already be stated by the function signature. However, it is difficult for the typechecker to figure out what type `x` might be. This is solved by typechecking `x` against *each* type that might occur (directly or indirectly) as subtree in `T`. If exactly one type matches, this typing is chosen. If not, an explicit typing is demanded.

This approach only works for complex expressions. Often, the programmer only wishes to capture the first occurrence of a certain syntactic form, which can be written as `e[(b:bool)]`. In order to save the programmer this boilerplate, the typechecker attempts to discover a syntactic form name in the variable type. If this name is found (as prefix), it will be inherently typed. In other words `e[bool]` is equivalent to `e[(bool:bool)]`.

```

1  typecheckExpression(expr,  $\gamma$ , T):
2      case expr of:
3          variable v:      return v:T
4          sequence es:
5              # includes lone string literals, sequence of one
6              possible_t typings = []
7              for choice_sequence in T.getChoices():
8                  if es.length != choice_sequence.length:
9                      continue
10                 try:
11                     typed_sequence = []
12                     for e, t in zip(es, choice_sequence):
13                         if e is literal && t is literal:
14                             if e != t then:
15                                 error "Inconsistent application"
16                             else typed_sequence += e
17                         else:
18                             typed_sequence += typecheckExpression(e,  $\gamma$ , t)
19                     possible_t typings += typed_sequence
20                 catch:
21                     # this doesn't match. Let's try the next choice...
22             if possible_t typings == []:
23                 error
24                 "Could not match $expr against"
25                 "any choice of the corresponding syntactic form"
26             if possible_t typings.length() > 1:
27                 error
28                 "Multiple possible typings for $expr."
29                 "Add an explicit type annotation"
30             return possible_t typings[0]
31     function f(x1, x2, ...):
32         (T1, T2, ..., RT) <-  $\gamma$ [f]      # Lookup type of f
33         x1' = typecheckExpression(x1)
34         x2' = typecheckExpression(x2)
35         if RT <: T:
36             return f(x1', x2', ...) : RT
37         else:
38             error
39             "Function $f does not have the desired type"
40     type annotation (e:TA):
41         if !(TA <: T):
42             error
43             "The typing annotation is too broad"
44             "or can never occur"
45         return typecheckExpression(e,  $\gamma$ , TA)
46     evaluation context e[x] with x a variable:
47         ts = T.occuringSubtypes().filter(x.isPrefixOf)
48         # occuringSubtypes are sorted on namelength
49         # the first match is the best match
50         t = ts[0]
51         typecheckExpression(e[(x:t)])
52     evaluation context e[x]:
53         ts = T.occuringSubtypes()
54         possible_t typings = []
55         for t in ts:
56             try{
57                 possible_t typings += typecheckExpression(x,  $\gamma$ , t)
58             }catch():
59                 # This doesn't match. Let's try the next one
60             if possible_t typings == []:
61                 error
62                 "Could not match $x against"
63                 "any possible embedded syntactic form"
64             if possible_t typings.length() > 1:
65                 error
66                 "Multiple possible typings for $x."
67                 "Add an explicit type annotation"
68             return possible_t typings[0]
69

```

Figure 1.5: The typechecking algorithm for meta-expressions and patterns


```

1  # Checks a clause for unknown or incompatible type variables
2  checkClause(pattern1, pattern2, ... , expr):
3      # search all the patterns for variables and their type
4      assign1 = pattern1.assignedVars()
5      assign2 = pattern2.assignedVars()
6      ...
7
8      assignE = expr.assignedVars()
9
10     assigns = merge(assign1, assign2, ...)
11
12     for variable_name in assignE.keys():
13         if !assigns.contains(variable_name):
14             error "Variable not defined"
15         TUsage = assignE.get(variable_name)
16         TDecl = assigns.get(variable_name)
17         if !TUsage.isSubtypeOf(TDecl):
18             error "Incompatible types"
19
20 merge(assign1, assign2, ...):
21     assign = {}
22     for variable_name in assign1.keys() + assign2.keys() + ... :
23         # assign.get(T) return Top for an unknown type
24         type = assign1.get(variable_name)
25              $\cap$  assign2.get(variable_name)
26              $\cap$  ...
27         if type is  $\epsilon$ :
28             error "Incompatible types while merging assignments"
29         assign.put(variable_name, type)
30     return assign

```

Figure 1.6: Merging of variable assignment stores and consistent variable usage checks

1.4 Interpretation of metafunctions

A function interpreter is builtin in ALGT, to evaluate the declared metafunctions. Just like any programming language, functions can be partial and fail. Failure can occur on two occasions:

- The builtin function `!error` is called
- A pattern match fails

1.4.1 Pattern matching

The input arguments of a function are matched against patterns for starters, simultaneously dispatching input cases and building a variable store. The pattern matcher is defined by giving behaviour for each possible pattern, as can be seen in figure 1.7.

The recursive nature of the pattern matching is visible in the case of `mSeq` (line 6) where the input `parsetree` is broken down in pieces and analyzed further.

Another interesting detail is the function case (line 27): the function is evaluated using `evaluate func`. If the function fails and gives an error message, then the pattern match will fail automatically and control flow moves to the next clause in the function, providing a rudimentary error recovery procedure.

1.4.2 Parsetree construction

Given a variable store, an expression can be evaluated easily. Mirroring `patternMatch`, `evaluate` gives a parsetree for each expression as can be seen in figure 1.8.

```

1 patternMatch :: Expression -> ParseTree -> Either String VariableAssignments
2 patternMatch (MVar v) expr
3     = return $ M.singleton v (expr, Nothing)
4
5 patternMatch (MParseTree (MLiteral _ _ s1)) (MLiteral _ _ s2)
6     | s1 == s2           = return M.empty
7     | otherwise          = Left $ "Not the same literal"
8
9 patternMatch (MParseTree (MInt _ _ s1)) (MInt _ _ s2)
10    | s1 == s2           = return M.empty
11    | otherwise          = Left $ "Not the same int"
12
13 patternMatch (MParseTree (PtSeq _ mi pts)) pt
14    = patternMatch (MSeq mi (pts |> MParseTree)) pt
15
16 patternMatch s1@(MSeq _ seq1) s2@(PtSeq _ _ seq2)
17    | length seq1 /= length seq2
18    = Left $ "Sequence lengths are not the same"
19    | otherwise
20    = zip seq1 seq2 |> uncurry (patternMatch )
21    >>= foldM mergeVars M.empty
22
23 patternMatch (MAscription as expr') expr
24    | alwaysIsA (typeof expr) as
25    = patternMatch expr' expr
26    | otherwise
27    = Left $ show expr ++ " is not a " ++ show as
28
29 patternMatch func@MCall{} arg
30    = do pt <- evaluate func
31        unless (pt == arg) $ Left $
32            "Function result does not equal the given argument"
33        return M.empty
34
35 patternMatch ctx _ pat expr
36    = Left $ "FT: Could not pattern match"

```

Figure 1.7: The pattern matching function. The code is edited for clarity, e.g. omitting a dictionary with known functions and the searching behaviour from evaluation contexts.

```

1 evaluate      :: VariableAssignments -> Expression -> Either String ParseTree
2 evaluate vars (MCall _ nm False args)
3   | nm `M.member` knownFunctions
4     = do      let func      = knownFunctions ! nm
5               args'      <- args |> evaluate ctx
6               applyFunc ctx (nm, func) args'
7   | otherwise
8     = evalErr ctx $ "unknown function: "++nm
9
10 evaluate vars (MVar _ nm)
11   | nm `M.member` vars
12     = return $ fst $ vars ! nm
13   | otherwise
14     = Left $ "Unkown variable"
15
16 evaluate vars (MSeq tp vals)
17   = do      vals' <- vals |> evaluate vars
18           return $ PtSeq vals'
19 evaluate vars (MParseTree pt)
20   = return pt
21 evaluate vars (MAscription tn expr)
22   = evaluate ctx expr

```

Figure 1.8: The parsetree construction, based on a meta-expression. The code is edited for clarity, e.g. omitting a dictionary with known functions, omitting evaluation contexts, ...

1.5 Natural Deduction

Building on the pattern matching blocks of function interpretation, natural deduction takes the evaluation onto a next level.

1.5.1 Representation

1.5.2 Typechecking

1.5.3 Interpretation

1.6 Correctness proof for STFL

For completeness, we include a proof for preservation and progress of our definition of STFL.